

Outcome Monitors: Recovering Silent Tool Failures in LLM Agents

Anonymous submission

Abstract

Modern artificial intelligence (AI) systems are powered by foundation models that increasingly act through external tools. When a tool call fails silently, the agent's downstream reasoning proceeds on stale or fabricated state. This paper presents Outcome Monitors, a lightweight observability layer that detects unacknowledged tool failures from the conversation trace alone. This abstract is placeholder text demonstrating the simple corner-logo variant on the standard AAAI template.

1 Introduction

Foundation models are general models of language, vision, speech, and/or other modalities. This body text exists only to show the standard AAAI layout with nothing changed except the corner logo. Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat.

Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

2 Related Work

Sed ut perspiciatis unde omnis iste natus error sit voluptatem accusantium doloremque laudantium, totam rem aperiam, eaque ipsa quae ab illo inventore veritatis et quasi architecto beatae vitae dicta sunt explicabo.