



Outcome Monitors: Recovering Silent Tool Failures in LLM Agents

Sugam Author¹, Rabab Abdelfattah¹

¹AIMS Lab, University of Southern Mississippi

Modern artificial intelligence (AI) systems are powered by foundation models that increasingly act through external tools. When a tool call fails silently, the agent's downstream reasoning proceeds on stale or fabricated state. This paper presents Outcome Monitors, a lightweight observability layer that detects unacknowledged tool failures from the conversation trace alone. This abstract is placeholder text demonstrating the AIMS Lab header on the ACL two-column template.

Date: August 13, 2026

Website: <https://aimslab.usm.edu>

1 Introduction

Foundation models are general models of language, vision, speech, and/or other modalities. This body text exists only to show the two-column ACL layout continuing below the full-width AIMS header box. Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat.

Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

2 Related Work

Sed ut perspiciatis unde omnis iste natus error sit voluptatem accusantium doloremque laudantium, totam rem aperiam, eaque ipsa quae ab illo inventore veritatis et quasi architecto beatae vitae dicta sunt explicabo.